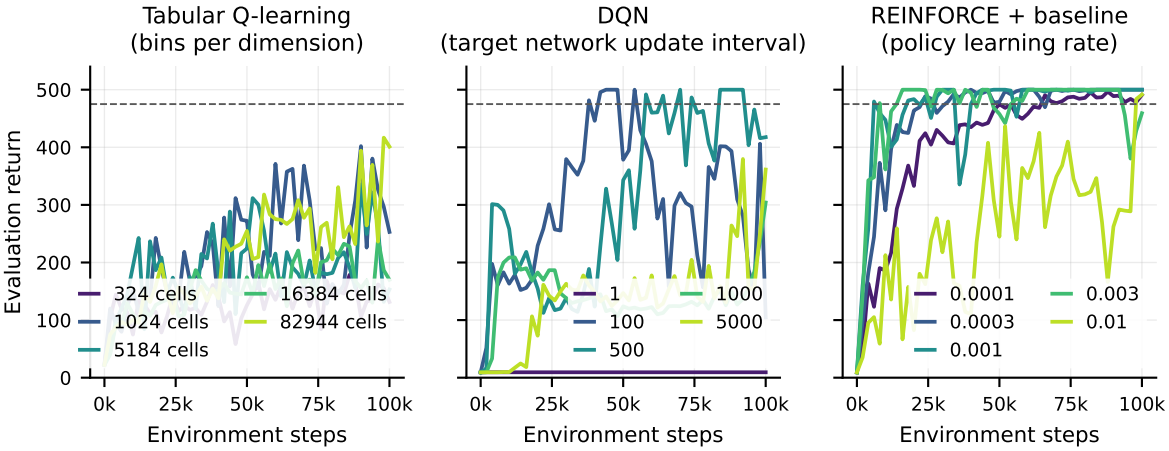
Tabular Q-learning
(bins per dimension)
DQN
(target network update interval)
REINFORCE + baseline
(policy learning rate)
Evaluation return
0
100
200
300
400
500
0k
25k
50k
75k
100k
Environment steps
324 cells
1024 cells
5184 cells
16384 cells
82944 cells
1
100
500
1000
5000
0.0001
0.0003
0.001
0.003
0.01